

more theory:

- recap: the inner product and its properties
- logistic sigmoid to map unbounded outputs onto  $[0, 1]$  with an interpretation as probability
- cross-entropy loss
- connection between the cross-entropy loss and the principle of maximum likelihood

- ① Important Intermezzo: the inner product
- ② Classification by logistic regression
- ③ Binary cross-entropy and Maximum Likelihood
- ④ Softmax and Cross-entropy loss for multiple classes

$$u \cdot v = \sum_{k=0}^{d-1} u_d v_d \in \mathbb{R}$$

has the following properties:

- ⊕ maps two real vectors  $u, v$  onto a real number  $u \cdot v$
- ⊕ linear in the first argument ( $u$ )

$$\begin{aligned} a \in \mathbb{R}, \quad (au) \cdot v &= a(u \cdot v) \\ (u^{\{1\}} + u^{\{2\}}) \cdot v &= u^{\{1\}} \cdot v + u^{\{2\}} \cdot v \end{aligned}$$

or in short:

$$(a_1 u^{\{1\}} + a_2 u^{\{2\}}) \cdot v = a_1(u^{\{1\}} \cdot v) + a_2(u^{\{2\}} \cdot v)$$

- ⊕ linear in the second argument ( $v$ )

$$u \cdot v = \sum_{k=0}^{d-1} u_d v_d \in \mathbb{R}$$

has the following properties:

$$\text{Symmetry: } u \cdot v = v \cdot u$$

$$v \neq \mathbf{0} \Rightarrow v \cdot v > 0$$

$$\mathbf{0} \cdot u = 0$$

where  $\mathbf{0}$  is the Null vector. For example in  $\mathbb{R}^3$  this is the element  $(0, 0, 0)$

$$u \cdot v = \sum_{k=0}^{d-1} u_d v_d \in \mathbb{R}$$

has the following properties:

- it defines a norm  $\|v\|$  is a norm, that is a notion of the length of a vector  $v$ :

$$v \cdot v = \|v\|^2$$

$$u \cdot v = \sum_{k=0}^{d-1} u_d v_d \in \mathbb{R}$$

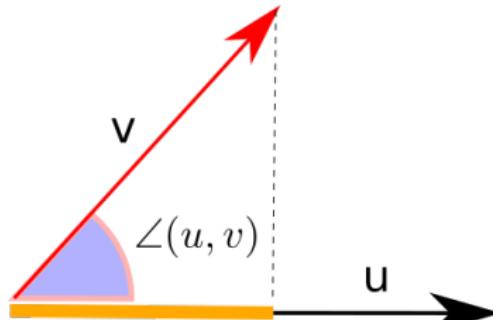
has the following properties:

- it defines an angle between two vectors:

$$\frac{u \cdot v}{(u \cdot u)^{1/2}(v \cdot v)^{1/2}} = \cos(\angle(u, v))$$

- the angle can be measured in any dimensions
- in higher dimensions, the angle is measured in the 2-dim plane spanned by  $u, v$ :

$$L(u, v) = \{a_0 u + a_1 v, a_0 \in \mathbb{R}, a_1 \in \mathbb{R}\}$$



has length equal to  
 $\|v\|_2 \cos(\angle(u, v))$

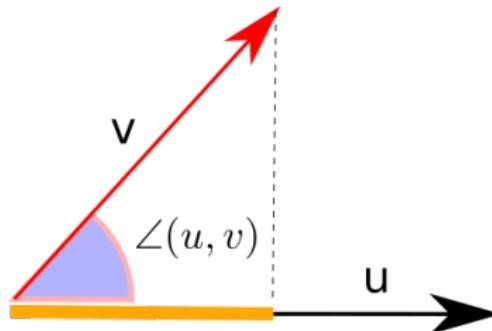
$$u \cdot v = \sum_{k=0}^{d-1} u_d v_d \in \mathbb{R}$$

- $u \cdot v$  defines an angle between two vectors:

$$\frac{u \cdot v}{(u \cdot u)^{1/2}(v \cdot v)^{1/2}} = \cos(\angle(u, v))$$

- for two vectors  $u, v$  of unit length ( $\|u\|_2 = 1$ ) the inner product

- lies in  $[-1, +1]$
- $u \cdot v = 1$  if  $u = v$
- gets close to 1 if their angle is close to zero,
- gets close to 0 if their angle is close to  $\pi/2 \sim 90$  deg,
- $u \cdot v = -1$  if  $u = -v$

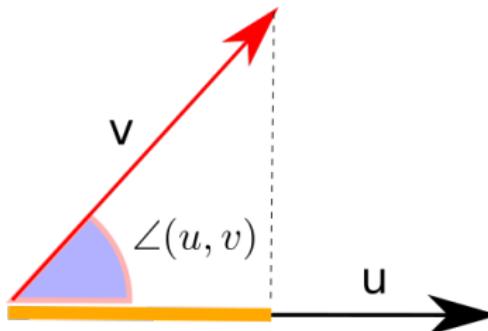


has length equal to  
 $\|v\|_2 \cos(\angle(u, v))$

$$u \cdot v = \sum_{k=0}^{d-1} u_d v_d \in \mathbb{R}$$

- $u \cdot v$  defines an angle between two vectors:

$$\frac{u \cdot v}{(u \cdot u)^{1/2}(v \cdot v)^{1/2}} = \cos(\angle(u, v))$$



has length equal to  
 $\|v\|_2 \cos(\angle(u, v))$

### Interpretation of the inner product

for two vectors  $u, v$  of unit length ( $\|u\|_2 = 1$ ) the inner product computes a similarity measure between  $u$  and  $v$  based on their angle!

### Interpretation of the inner product

for two vectors  $u, v$  of unit length ( $\|u\|_2 = 1$ ) the inner product computes a similarity measure between  $u$  and  $v$  based on their angle!

next step:

- now we can use  $u \cdot v$  to define a simple classifier:

$$\begin{aligned}f(x) &= w \cdot x + b \\s(x) &= \text{sign}(f(x)) \in \{-1, +1\}\end{aligned}$$

Mechanism:

- $f(x)$  is large if the angle  $\angle(w, x)$  is close to zero,
- it assigns large values to  $x$  with  $\angle(w, x)$  close to zero,

Simplest neural network:

- $x = (x_0, x_1, \dots, x_{d-1})^\top \in \mathbb{R}^d$  - input vector.
- $f(x)$  output of the only weight layer
- with weight vector  $w = (w_0, w_1, \dots, w_{d-1}, b) \in \mathbb{R}^d$
- $y = \text{sign}(z)$  - activation function on top of  $f(x)$

- ① Important Intermezzo: the inner product
- ② Classification by logistic regression
- ③ Binary cross-entropy and Maximum Likelihood
- ④ Softmax and Cross-entropy loss for multiple classes

Next steps: There is lots to derive from this simple case  $f(x) = w \cdot x + b$

- derive an activation function with nice properties (→ logistic sigmoid)
- derive a loss for this case (→ Binary cross entropy loss)
- extend this to more than 2 classes as outputs (→ softmax)
- derive a loss function for more than 2 classes as outputs (→ Cross entropy loss)

Goal of classification: For every input sample  $x \in \mathcal{X}$ , correctly predict which class  $y$  it belongs.

- ⊕ For 2-class classification,  $y \in \{-1, +1\}$  or  $y \in \{0, 1\}$ .

First attempt: Apply a linear mapping and classify according to the sign of the output:

$$f(x) = w \cdot x + b, \quad s(x) = \text{sign}(f(x)) \in \{-1, +1\}$$

While  $f(x)$  has unbounded values,  $s(x)$  is either  $-1$  or  $1$ , but:

- ⊕ if  $f(x) \approx 0$ , we should be uncertain about the prediction.
- ⊕ if  $f(x) \gg 0$ , we should be confident about the prediction  $1$ .
- ⊕ if  $f(x) \ll 0$ , we should be confident about the prediction  $-1$ .

Goal: encode this uncertainty

Goal of classification: For every input sample  $x \in \mathcal{X}$ , correctly predict which class  $y$  it belongs.

- ⊕ For 2-class classification,  $y \in \{-1, +1\}$  or  $y \in \{0, 1\}$ .

First attempt: Apply a linear mapping and classify according to the sign of the output:

$$f(x) = w \cdot x + b, \quad s(x) = \text{sign}(f(x)) \in \{-1, +1\}$$

While  $f(x)$  has unbounded values,  $s(x)$  is either  $-1$  or  $1$ , but:

- ⊕ if  $f(x) \approx 0$ , we should be uncertain about the prediction.
- ⊕ if  $f(x) \gg 0$ , we should be confident about the prediction  $1$ .
- ⊕ if  $f(x) \ll 0$ , we should be confident about the prediction  $-1$ .

Goal: encode this uncertainty

Goal of classification: For every input sample  $x \in \mathcal{X}$ , correctly predict which class  $y$  it belongs.

- For 2-class classification,  $y \in \{-1, +1\}$  or  $y \in \{0, 1\}$ .

First attempt: Apply a linear mapping and classify according to the sign of the output:

$$f(x) = w \cdot x + b, \quad s(x) = \text{sign}(f(x)) \in \{-1, +1\}$$

While  $f(x)$  has unbounded values,  $s(x)$  is either  $-1$  or  $1$ , but:

- if  $f(x) \approx 0$ , we should be uncertain about the prediction.
- if  $f(x) \gg 0$ , we should be confident about the prediction  $1$ .
- if  $f(x) \ll 0$ , we should be confident about the prediction  $-1$ .

Goal: encode this uncertainty

Goal of classification: For every input sample  $x \in \mathcal{X}$ , correctly predict which class  $y$  it belongs.

- For 2-class classification,  $y \in \{-1, +1\}$  or  $y \in \{0, 1\}$ .

First attempt: Apply a linear mapping and classify according to the sign of the output:

$$f(x) = w \cdot x + b, \quad s(x) = \text{sign}(f(x)) \in \{-1, +1\}$$

While  $f(x)$  has unbounded values,  $s(x)$  is either  $-1$  or  $1$ , but:

- if  $f(x) \approx 0$ , we should be uncertain about the prediction.
- if  $f(x) \gg 0$ , we should be confident about the prediction  $1$ .
- if  $f(x) \ll 0$ , we should be confident about the prediction  $-1$ .

Goal: encode this uncertainty

Goal of classification: For every input sample  $x \in \mathcal{X}$ , correctly predict which class  $y$  it belongs.

- For 2-class classification,  $y \in \{-1, +1\}$  or  $y \in \{0, 1\}$ .

First attempt: Apply a linear mapping and classify according to the sign of the output:

$$f(x) = w \cdot x + b, \quad s(x) = \text{sign}(f(x)) \in \{-1, +1\}$$

While  $f(x)$  has unbounded values,  $s(x)$  is either  $-1$  or  $1$ , but:

- if  $f(x) \approx 0$ , we should be uncertain about the prediction.
- if  $f(x) \gg 0$ , we should be confident about the prediction  $1$ .
- if  $f(x) \ll 0$ , we should be confident about the prediction  $-1$ .

Goal: encode this uncertainty

How can we encode the uncertainty into a mapping from  $(-\infty, +\infty)$  onto  $[0, 1]$ ? We would like to map using a function  $s(\cdot)$ :

- $f(x) \approx 0$  to  $s(f(x)) \approx 0.5$ .
- $f(x) \gg 0$  to  $s(f(x)) \approx 1$ , in particular let  $\lim_{u \rightarrow +\infty} s(u) = 1$
- $f(x) \ll 0$  to  $s(f(x)) \approx 0$ , in particular let  $\lim_{u \rightarrow -\infty} s(u) = 0$

This would allow us to interpret  $s(u)$  as a probability over inputs  $u$ .

There are many possible choices for the function  $s(u)$ .

How can we encode the uncertainty into a mapping from  $(-\infty, +\infty)$  onto  $[0, 1]$ ? We would like to map using a function  $s(\cdot)$ :

- $f(x) \approx 0$  to  $s(f(x)) \approx 0.5$ .
- $f(x) \gg 0$  to  $s(f(x)) \approx 1$ , in particular let  $\lim_{u \rightarrow +\infty} s(u) = 1$
- $f(x) \ll 0$  to  $s(f(x)) \approx 0$ , in particular let  $\lim_{u \rightarrow -\infty} s(u) = 0$

This would allow us to interpret  $s(u)$  as a probability over inputs  $u$ .

There are many possible choices for the function  $s(u)$ .

How can we encode the uncertainty into a mapping from  $(-\infty, +\infty)$  onto  $[0, 1]$ ? We would like to map using a function  $s(\cdot)$ :

- $f(x) \approx 0$  to  $s(f(x)) \approx 0.5$ .
- $f(x) \gg 0$  to  $s(f(x)) \approx 1$ , in particular let  $\lim_{u \rightarrow +\infty} s(u) = 1$
- $f(x) \ll 0$  to  $s(f(x)) \approx 0$ , in particular let  $\lim_{u \rightarrow -\infty} s(u) = 0$

This would allow us to interpret  $s(u)$  as a probability over inputs  $u$ .

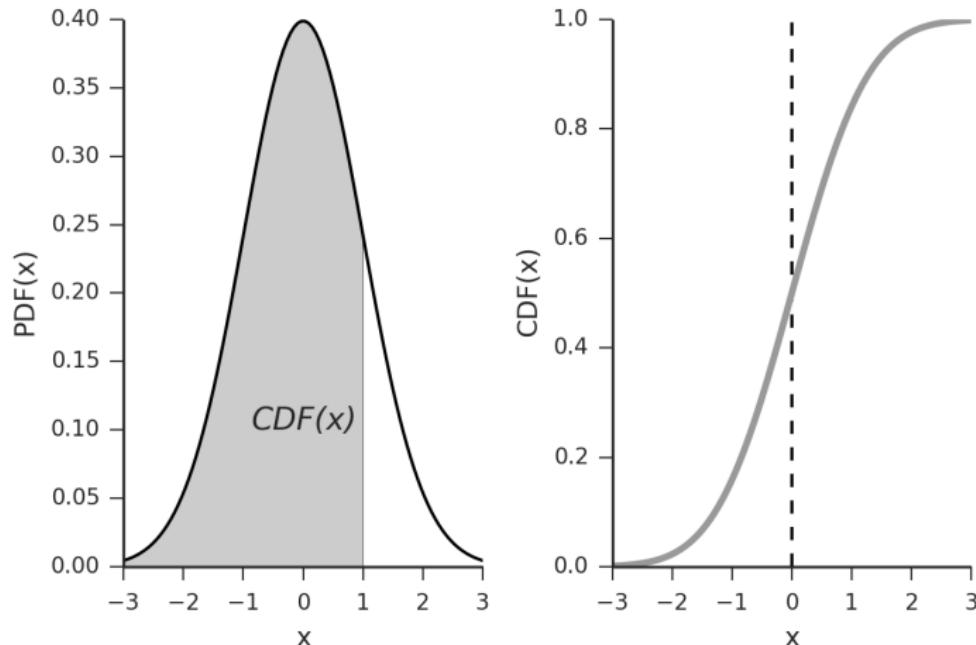
There are many possible choices for the function  $s(u)$ .

There are many possible choices for the function  $s(u)$ .



We could e.g. use the cumulative distribution function (CDF) of any probability distribution function (PDF) with a median of 0.

- The normal (Gaussian) distribution is used in probit (from **probability and unit**) regression.



We will use a simpler function called the logistic sigmoid function:

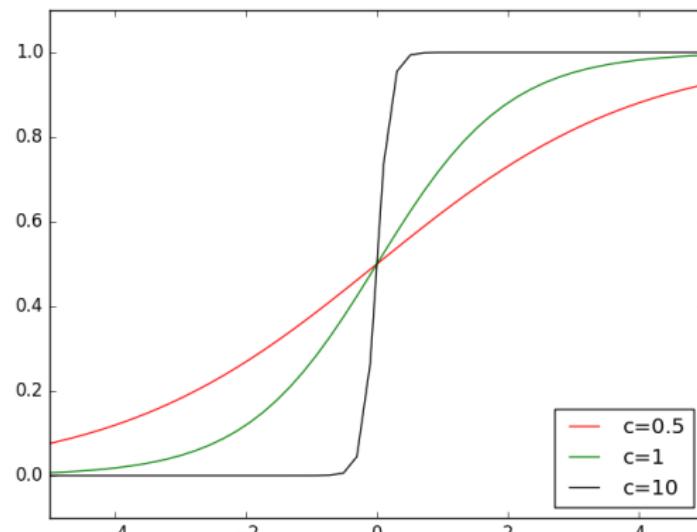
Definition: Logistic sigmoid function

$$s(u) = \frac{\exp(u)}{1 + \exp(u)} = \frac{1}{\exp(-u) + 1} \frac{\exp(u)}{\exp(u)} = \frac{1}{\exp(-u) + 1}$$

How does the logistic sigmoid function look like?

Let's plot it for different scaling factors  $c > 0$ :

$$s(cu) = \frac{\exp(cu)}{1 + \exp(cu)} = \frac{1}{\exp(-cu) + 1}$$



Definition: Logistic sigmoid function

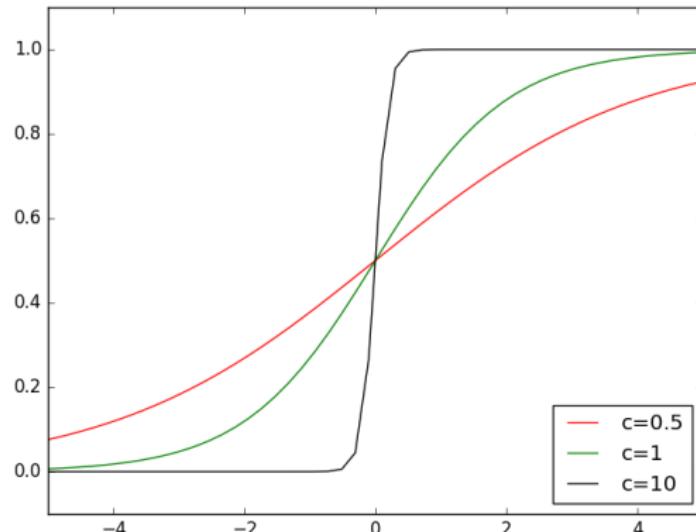
$$s(u) = \frac{\exp(u)}{1 + \exp(u)} = \frac{1}{\exp(-u) + 1} \frac{\exp(u)}{\exp(u)} = \frac{1}{\exp(-u) + 1}$$

Note that:

$$s(0) = \frac{1}{\exp(-0) + 1} = \frac{1}{1 + 1} = 0.5$$

$$\lim_{u \rightarrow \infty} s(u) = \lim_{u \rightarrow \infty} \frac{1}{\exp(-u) + 1} = \frac{1}{0 + 1} = 1$$

$$\lim_{u \rightarrow -\infty} s(u) = \lim_{u \rightarrow -\infty} \frac{\exp(u)}{1 + \exp(u)} = \frac{0}{1 + 0} = 0$$



### Definition: Logistic regression model

Assume we have an affine mapping  $f_{w,b}(x) = w \cdot x + b$ . Plugging it into the logistic sigmoid function  $s(u)$  provides a logistic regression model:

$$s(f_{w,b}(x)) = \frac{\exp(w \cdot x + b)}{1 + \exp(w \cdot x + b)} = \frac{1}{\exp(-w \cdot x - b) + 1}$$

For 2-class classification problems, the output  $s(f_{w,b}(x)) \in [0, 1]$  is the predicted probability that sample  $x$  has class label  $y = 1$ , that is,  $P(Y = 1 | X = x)$ .

## Interpretation of Logistic regression model outputs

$$s(f_{w,b}(x)) = \frac{\exp(w \cdot x + b)}{1 + \exp(w \cdot x + b)} = \frac{1}{\exp(-w \cdot x - b) + 1}$$

is a model for  $P(Y = 1|X = x)$  in the classification prediction.

Indeed, for every input  $x$  we have that  $s(f_{w,b}(x)) \in (0, 1)$ . So it can be considered as a probability value.

Since we have only two classes, it must hold:

$$\begin{aligned} P(Y = 1|X = x) + P(Y = 0|X = x) &= 1 \\ P(Y = 0|X = x) &= 1 - P(Y = 1|X = x) = 1 - s(f_{w,b}(x)) \end{aligned}$$

and we can use  $s(f_{w,b}(x))$  to express the probability for the other class  $P(Y = 0|X = x)$ .

next step:

- derive a loss function which can be used for the logistic regression prediction model!

### Binary cross entropy loss for a single output

Let us consider classification with 2 classes with labels  $\{0, 1\}$ , and let  $s(x)$  be a model for  $P(Y = 1|X = x)$ . Then the binary cross entropy loss for a pair of prediction and label  $(s(x), y)$  is given as

$$e(x, y) = -y \ln(s(x)) - (1 - y) \ln(1 - s(x))$$

This looks complicated but it has a simple interpretation:

- $y = 1 \Rightarrow e = -\ln(s(x))$  ...! the neg-log-probability  $P(Y = 1|X = x)$  of the ground truth class  $y = 1$
- $y = 0 \Rightarrow e = -\ln(1 - s(x))$  ...! the neg-log-probability  $P(Y = 0|X = x)$  of the ground truth class  $y = 0$

### Binary cross entropy loss for a single output

Let us consider classification with 2 classes with labels  $\{0, 1\}$ , and let  $s(x)$  be a model for  $P(Y = 1|X = x)$ . Then the binary cross entropy loss for a pair of prediction and label  $(s(x), y)$  is given as

$$e(x, y) = -\ln(P(Y = 1|X = x))1[y == 1] - \ln(P(Y = 0|X = x))1[y == 0]$$

### Binary cross entropy loss for a single output

Let us consider classification with 2 classes with labels  $\{0, 1\}$ , and let  $s(x)$  be a model for  $P(Y = 1|X = x)$ . Then the binary cross entropy loss for a pair of prediction and label  $(s(x), y)$  is given as

$$e = -y \ln(s(x)) - (1 - y) \ln(1 - s(x))$$

This is the neg-logarithm of the probability  $P(Y = y|X = x)$  for the ground truth label class  $y$  where  $s(x) = P(Y = 1|X = x)$  – if one uses 0-1-labels.  
(for a given pair  $(x, y)$  of feature  $x$  and its label  $y$ )

Next: Why is the neg-logarithm of the probability of the ground-truth class a good loss function?

Next: Why is the neg-logarithm of the probability of the ground-truth class a good loss function?

Desirable properties for a loss function:

- it should have low loss if the prediction is close to the ground truth
- it should have high loss if the prediction is far from the ground truth
- if one wants to use gradient-based optimization, it should be almost everywhere differentiable

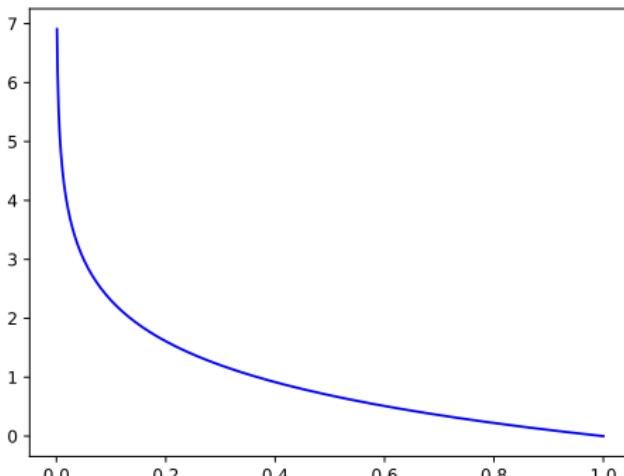
# Why is the neg-logarithm of the probability of the ground-truth class a good loss function?

| 26

Suppose  $y = 1$ . Then the optimal prediction is  $P(Y = 1|X = x) = 1$ .

- If we had the optimal prediction is  $P(Y = 1|X = x) = 1$ , then the loss is  $-\ln(1) = 0$ . Zero loss
- If we had the least desirable prediction  $P(Y = 1|X = x) = 0$ , then the loss is  $\lim_{x \rightarrow 0, x > 0} -\ln(x) = \infty$ .

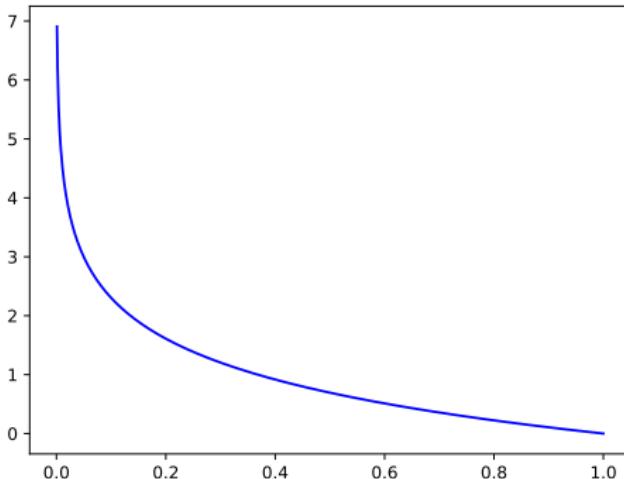
The neg logarithm looks like this:



# Why is the neg-logarithm of the probability of the ground-truth class a good loss function?

| 27

The neg logarithm looks like this:



- So if  $y = 1$ , then too low predicted probabilities  $s(x)$  close to zero get a large loss.
- Similarly if  $y = 0$ , then too high probabilities  $s(x) \approx 1$  close to one get a large loss (because then  $1 - s(x) \approx 0$  is close to zero again, thus its neg log will be high)

Takeaway:

- the neg log probability prediction for the ground truth class is a reasonable loss function for classification problems.
- we have a classification model, which returns probabilities
- we have a loss to train it (not said how to use it)

Next:

- show that this loss has a derivation from a principle.

- ① Important Intermezzo: the inner product
- ② Classification by logistic regression
- ③ **Binary cross-entropy and Maximum Likelihood**
- ④ Softmax and Cross-entropy loss for multiple classes

Next:

- show that this loss has a derivation from a principle.

Steps:

- (1) define a probability model  $P(Y|X = x)$  for the prediction on a single input sample  $x$
- (2) define a probability model  $P(Y_0, Y_1, Y_2, \dots | X_0 = x_0, X_1 = x_1, X_2 = x_2, \dots)$  for predictions on a set of samples  $x_0, x_1, x_2, \dots$
- (3) employ the principle of maximum likelihood to select parameters for the probability model
- (4) use a log-space transformation

(a probability model  $P(Y|X = x)$ )

- The logistic output  $s(f_{w,b}(x)) \in [0, 1]$  can be interpreted as the probability  $P(\{Y = 1|X = x\})$  of observing label  $Y = 1$ , predicted by the model.
- then:

$$P(Y = 0|X = x) = 1 - P(Y = 1|X = x) = 1 - s(f_{w,b}(x))$$

$$y \in \{0, 1\} : P(Y = y|X = x) = sy + (1 - s)(1 - y)$$

$$y \in \{0, 1\} : P(Y = y|X = x) = s^y(1 - s)^{1-y}$$

both are valid expressions for  $P(Y = y|X = x)$

- we will use  $y \in \{0, 1\} : P(Y = y|X = x) = s^y(1 - s)^{1-y}$  ( compat. w. log-transform)

- define a model for  $P(Y_0, Y_1, Y_2, \dots | X_0 = x_0, X_1 = x_1, X_2 = x_2, \dots)$  by multiplying all probabilities  $P(\{Y_k | X_k = x_k\})$  together:

$$\begin{aligned} P(Y_0, Y_1, \dots, Y_{n-1} | X_0 = x_0, X_1 = x_1, \dots, X_{n-1} = x_{n-1}) &= \\ P(\{Y_0 | X_0 = x_0\})P(\{Y_1 | X_1 = x_1\}) \cdot \dots \cdot P(\{Y_{n-1} | X_{n-1} = x_{n-1}\}) \\ &= \prod_{k=0}^{n-1} P(\{Y_k | X_k = x_k\}) \end{aligned}$$

- why does this make sense?
  - note: The set of all values for  $(Y_0, Y_1, \dots, Y_{n-1})$  are all sequences  $(\underbrace{11010\dots1}_{len=n})$  of length  $n$ .
  - example values for  $n = 3$  are  $(0, 0, 0), (1, 0, 1), (0, 1, 1)$
  - Expressed by math:  $\{0, 1\}^n$

$$\begin{aligned}
 P(Y_0, Y_1, \dots, Y_{n-1} | X_0 = x_0, X_1 = x_1, \dots, X_{n-1} = x_{n-1}) &= \\
 &= \prod_{k=0}^{n-1} P(\{Y_k | X_k = x_k\})
 \end{aligned}$$

- >this is a probability over all possible values of  $Y_0, Y_1, \dots, Y_{n-1}$ :

- $P(\{Y_k | X_k = x_k\}) \in [0, 1] \Rightarrow 0 \leq \prod_{k=0}^{n-1} P(\{Y_k | X_k = x_k\}) \leq 1$
- it sums up to = 1 over all possible values:

$$\sum_{Y_0 \in \{0,1\}, Y_1 \in \{0,1\}, \dots, Y_{n-1} \in \{0,1\}} \prod_{k=0}^{n-1} P(\{Y_k | X_k = x_k\})$$

$$= \sum_{Y_0 \in \{0,1\}} \sum_{Y_1 \in \{0,1\}} \dots \sum_{Y_{n-1} \in \{0,1\}} \prod_{k=0}^{n-1} P(\{Y_k | X_k = x_k\}) = 1$$

$$\begin{aligned} & \sum_{Y_0 \in \{0,1\}, Y_1 \in \{0,1\}, \dots, Y_{n-1} \in \{0,1\}} \prod_{k=0}^{n-1} P(\{Y_k | X_k = x_k\}) \\ &= \sum_{Y_0 \in \{0,1\}} \sum_{Y_1 \in \{0,1\}} \dots \sum_{Y_{n-1} \in \{0,1\}} \prod_{k=0}^{n-1} P(\{Y_k | X_k = x_k\}) \\ &= \sum_{Y_0 \in \{0,1\}} \sum_{Y_1 \in \{0,1\}} \dots \sum_{Y_{n-1} \in \{0,1\}} P(\{Y_0 | X_0 = x_0\}) \prod_{k=1}^{n-1} P(\{Y_k | X_k = x_k\}) \\ & P(Y_0 | \dots) \text{ is a constant for summations over } Y_1, Y_2, \dots, Y_{n-1} \end{aligned}$$

The trick of pulling out a constant:

$$\begin{aligned} \sum_{Y_1} \sum_{Y_2} c(Y_0) t(Y_1, Y_2) &= c(Y_0) \sum_{Y_1} \sum_{Y_2} t(Y_1, Y_2) \\ \Rightarrow \sum_{Y_0} \sum_{Y_1} \sum_{Y_2} c(Y_0) t(Y_1, Y_2) &= \sum_{Y_0} c(Y_0) \sum_{Y_1} \sum_{Y_2} t(Y_1, Y_2) \end{aligned}$$

$$\begin{aligned}
 & \sum_{Y_0 \in \{0,1\}, Y_1 \in \{0,1\}, \dots, Y_{n-1} \in \{0,1\}} \prod_{k=0}^{n-1} P(\{Y_k | X_k = x_k\}) \\
 &= \sum_{Y_0 \in \{0,1\}} \sum_{Y_1 \in \{0,1\}} \dots \sum_{Y_{n-1} \in \{0,1\}} P(\{Y_0 | X_0 = x_0\}) \prod_{k=1}^{n-1} P(\{Y_k | X_k = x_k\}) \\
 & \quad P(Y_0 | \dots) \text{ is a constant for summations over } Y_1, Y_2, \dots, Y_{n-1} \\
 &= \sum_{Y_0 \in \{0,1\}} P(\{Y_0 | X_0 = x_0\}) \sum_{Y_1 \in \{0,1\}} \dots \sum_{Y_{n-1} \in \{0,1\}} \prod_{k=1}^{n-1} P(\{Y_k | X_k = x_k\}) \\
 &= 1 \sum_{Y_1 \in \{0,1\}} \dots \sum_{Y_{n-1} \in \{0,1\}} \prod_{k=1}^{n-1} P(\{Y_k | X_k = x_k\})
 \end{aligned}$$

n-1 times iterate:  $= 1 \cdot \dots \cdot 1 = 1$

- Result: We have shown that

$$\begin{aligned} P(Y_0, Y_1, \dots, Y_{n-1} | X_0 = x_0, X_1 = x_1, \dots, X_{n-1} = x_{n-1}) &= \\ &= \prod_{k=0}^{n-1} P(\{Y_k | X_k = x_k\}) \end{aligned}$$

defines a probability distribution over the set  $\{0, 1\}^n$  of all values for the sequence  $(Y_0, Y_1, \dots, Y_{n-1})$

- using our logistic regression output  $s(f_{w,b}(x))$ , we have defined a probability model for observing values  $y_0, y_1, \dots, y_{n-1}$

$$P(Y_0, Y_1, \dots, Y_{n-1} | X_0 = x_0, X_1 = x_1, \dots, X_{n-1} = x_{n-1}, w, b)$$

- we would like to find suitable values for the model parameters  $w$  and  $b$ .
- the simplest thought is to choose  $(w, b)$  such that the probability to observe the data which we have  $y_0, \dots, y_{n-1} \in \{0, 1\}$  is maximized.
- This is called the principle of maximal likelihood  
(wähle die Parameter derart, dass die Plausibilität der Beobachtung der gegebenen Daten maximiert wird)

$$(w^*, b^*) = \operatorname{argmax}_{\{(w,b)\}} P(Y_0 = y_0, Y_1 = y_1, \dots, Y_{n-1} = y_{n-1} | w, b)$$

- using our logistic regression output  $s(f_{w,b}(x))$ , we have defined a probability model for observing values  $y_0, y_1, \dots, y_{n-1}$

$$P(Y_0, Y_1, \dots, Y_{n-1} | X_0 = x_0, X_1 = x_1, \dots, X_{n-1} = x_{n-1}, w, b)$$

- we would like to find suitable values for the model parameters  $w$  and  $b$ .
- the simplest thought is to choose  $(w, b)$  such that the probability to observe the data which we have  $y_0, \dots, y_{n-1} \in \{0, 1\}$  is maximized.
- This is called the principle of maximal likelihood  
(wähle die Parameter derart, dass die Plausibilität der Beobachtung der gegebenen Daten maximiert wird)

$$(w^*, b^*) = \operatorname{argmax}_{\{(w,b)\}} P(Y_0 = y_0, Y_1 = y_1, \dots, Y_{n-1} = y_{n-1} | w, b)$$

### Maximum likelihood principle

Given some data observations  $z_0, z_1, \dots, z_{n-1}$ , and a probability model  $P(Z_0 = z_0, Z_1 = z_1, \dots, Z_{n-1} = z_{n-1} | \theta)$  which depends on some parameter  $\theta$ , the Maximum likelihood principle says, that one can choose  $\theta$  such that it maximizes the probability of observing the given data:

$$\theta^* = \operatorname{argmax}_{\text{all possible } \theta} P(Z_0 = z_0, Z_1 = z_1, \dots, Z_{n-1} = z_{n-1} | \theta)$$

Choose the model parameters such that observing the given data samples becomes most likely.

One missing point:

- Why does it make sense to define a probability as a product of observations for single samples ?

$$\begin{aligned} P(Y_0, Y_1, \dots, Y_{n-1} | X_0 = x_0, X_1 = x_1, \dots, X_{n-1} = x_{n-1}) &= \\ &= \prod_{k=0}^{n-1} P(\{Y_k | X_k = x_k\}) \end{aligned}$$

This makes sense if one can assume statistical independence (to be exact: conditional independence of  $Y$  given that one knows the  $X$ )

Give an example of data collection, where statistical independence is not fully satisfied?



## Maximum likelihood principle with independence assumption

Given  $n$  vectors/data points  $z_0, \dots, z_{n-1}$ , a probability model  $p(z_i; \theta)$  and an assumption of independence of  $z_i$  given  $\theta$ , choose the parameters  $\theta$  such that the probability of observing the data is maximized:

$$\theta^* = \operatorname{argmax}_{\theta} p(z_0; \theta)p(z_1; \theta) \dots p(z_{n-1}; \theta)$$

$$\text{or } \theta^* = \operatorname{argmin}_{\theta} \sum_{i=0}^{n-1} -\ln(p(z_i; \theta))$$

Use as model for generating / simulating new data:  $p(x; \theta^*)$

Maximum likelihood with independence assumption:

$$\theta^* = \operatorname{argmax}_\theta p(z_0; \theta)p(z_1; \theta) \dots p(z_{n-1}; \theta)$$

for many distributions with exponentials it is easier to instead minimize the neg-log likelihood.

$$\begin{aligned}\theta^* &= \operatorname{argmin}_\theta - \ln \left[ p(z_0; \theta)p(z_1; \theta) \dots p(z_{n-1}; \theta) \right] \\ &= \operatorname{argmin}_\theta \left[ -\ln p(z_0; \theta) - \ln p(z_1; \theta) - \dots - \ln p(z_{n-1}; \theta) \right]\end{aligned}$$

This gives the same solution.

Next step: apply this:

$$\operatorname{argmin}_{\theta} \left[ -\ln p(z_0; \theta) - \ln p(z_1; \theta) - \dots - \ln p(z_{n-1}; \theta) \right]$$

- ⊕ plug in our model:  $z_i = (y_i, z_i)$ ,  $P(z_i; \theta) = P(Y_i = y_i | X_i = x_i, w, b)$
- ⊕ use  $P(Y_i = y_i | X_i = x_i, w, b) = s(x_i)^{y_i} (1 - s(x_i))^{1-y_i}$  and  
 $\ln(ab) = \ln(a) + \ln(b)$ ,  $\ln(a^c) = c \ln(a)$

$$\begin{aligned}-\ln P(Y_i = y_i | X_i = x_i, w, b) &= -\ln(s(x_i)^{y_i} (1 - s(x_i))^{1-y_i}) \\ &= y_i(-\ln(s(x_i))) + (1 - y_i)(-\ln(1 - s(x_i)))\end{aligned}$$

This is the binary cross entropy loss from above (compare!)

$$e = -y \ln(s(x)) - (1 - y) \ln(1 - s(x))$$

Take away:

### Binary cross entropy loss for a single output

Let us consider classification with 2 classes with labels  $\{0, 1\}$ , and let  $s(x)$  be a model for  $P(Y = 1|X = x)$ . Then the binary cross entropy loss for a pair of prediction and label  $(s(x), y)$  is given as

$$e = -y \ln(s(x)) - (1 - y) \ln(1 - s(x)) = \sum_{i \in \{0,1\}} -\ln P(Y = i|X = x) \mathbb{1}[i = y]$$

- This is the neg-logarithm of the probability  $P(Y = y|X = x)$  for the ground truth label class  $y$
- It can be derived from applying the maximum likelihood principle aiming at choosing parameters to maximize the probability to observe the ground truth labels  $y_i$

- ① Important Intermezzo: the inner product
- ② Classification by logistic regression
- ③ Binary cross-entropy and Maximum Likelihood
- ④ Softmax and Cross-entropy loss for multiple classes

- Suppose we have  $C$  classes:  $y \in \{0, \dots, C - 1\}$  and
- we have a function  $P(Y = k | X_i = x_i, \theta)$  which returns probabilities for each label  $k$
- the problem is **Multi-class classification**: the  $C$  classes are mutually exclusive.

### Multi-class classification

- class labels are mutually exclusive in the ground truth
- if the predictions modeled as probability  $P(Y = k | X_i = x_i, \theta)$ , this requires that

$$\sum_{k=0}^{C-1} P(Y = k | X_i = x_i, \theta) = 1$$

- note: one can model the absence of anything by a background class label

- Suppose we have  $C$  classes:  $y \in \{0, \dots, C - 1\}$  and
- we have a function  $P(Y = k | X_i = x_i, \theta)$  which returns probabilities for each label  $k$
- Lets represent the label  $y \in \{0, \dots, C - 1\}$  via a one-hot vector  $\tilde{y} \in \mathbb{R}^c$ :

$$\tilde{y}[k] = \begin{cases} 1 & \text{if } y = k \\ 0 & \text{else} \end{cases}$$
$$\tilde{y} = (0, \dots, 0, \underbrace{1}_{\text{at value of } y}, 0, \dots)^\top$$

Then  $P(Y_i = y_i | X_i = x_i, \theta) = \prod_{k=0}^{C-1} P(Y = k | X_i = x_i, \theta)^{\tilde{y}_i[k]}$

- Suppose we have  $C$  classes:  $y \in \{0, \dots, C - 1\}$  and
- we have a function  $P(Y = k | X_i = x_i, \theta)$  which returns probabilities for each label  $k$
- Lets represent the label  $y \in \{0, \dots, C - 1\}$  via a one-hot vector  $\tilde{y} \in \mathbb{R}^C$

Then  $P(Y_i = y_i | X_i = x_i, \theta) = \prod_{k=0}^{C-1} P(Y = k | X_i = x_i, \theta)^{\tilde{y}_i[k]}$

... simply because for all  $k$  where  $\tilde{y}_i[k] = 0$  we have

$$P(Y = k | X_i = x_i, \theta)^{\tilde{y}_i[k]} = P(Y = k | X_i = x_i, \theta)^0 = 1$$

... for the sole one  $k$  where  $\tilde{y}_i[k] = 1$  this returns:

$$\begin{aligned} P(Y = k | X_i = x_i, \theta)^{\tilde{y}_i[k]} &= P(Y = k | X_i = x_i, \theta)^1 = P(Y_i = y_i | X_i = x_i, \theta) \\ \Rightarrow \prod_{k=0}^{C-1} P(Y = k | X_i = x_i, \theta)^{\tilde{y}_i[k]} &= P(Y_i = y_i | X_i = x_i, \theta) \end{aligned}$$

We have have  $C$  classes:  $y \in \{0, \dots, C - 1\}$  and

$$\prod_{k=0}^{C-1} P(Y = k | X_i = x_i, \theta)^{\tilde{y}_i[k]} = P(Y_i = y_i | X_i = x_i, \theta)$$

now apply the same principle: we try to find  $\theta^*$  which is the minimizer of

$$\begin{aligned}\theta^* &= \operatorname{argmin}_{\theta} \sum_i -\ln p(z_i; \theta) \\ &= \operatorname{argmin}_{\theta} \sum_i -\ln P(Y_i = y_i | X_i = x_i, \theta) \\ &= \operatorname{argmin}_{\theta} \sum_i -\ln \prod_{k=0}^{C-1} P(Y = k | X_i = x_i, \theta)^{\tilde{y}_i[k]} \\ &= \operatorname{argmin}_{\theta} \sum_i \sum_{k=0}^{C-1} -\ln P(Y = k | X_i = x_i, \theta)^{\tilde{y}_i[k]} \\ &= \operatorname{argmin}_{\theta} \sum_i \sum_{k=0}^{C-1} -\tilde{y}_i[k] \ln P(Y = k | X_i = x_i, \theta)\end{aligned}$$

We have have  $C$  classes:  $y \in \{0, \dots, C - 1\}$  and

$$\prod_{k=0}^{C-1} P(Y = k | X_i = x_i, \theta)^{\tilde{y}_i[k]} = P(Y_i = y_i | X_i = x_i, \theta)$$

we try to find  $\theta^*$  which is the minimizer of

$$\theta^* = \operatorname{argmin}_{\theta} \sum_{\text{samples } i} \sum_{k=0}^{C-1} -\tilde{y}_i[k] \ln P(Y = k | X_i = x_i, \theta)$$

### Multi-class cross entropy loss

Let  $\tilde{y}_i$  be the  $C$ -dimensional one-hot vector encoding labels in  $\{0, \dots, C - 1\}$ , and let  $P(Y = k | X_i = x_i, \theta)$  the probability for class  $k$  given input sample  $x_i$ , then the cross-entropy loss for one sample  $(x_i, y_i)$  is defined as:

$$\sum_{k=0}^{C-1} -\tilde{y}_i[k] \ln P(Y = k | X_i = x_i, \theta)$$

### Multi-class cross entropy loss

Let  $\tilde{y}_i$  be the  $C$ -dimensional one-hot vector encoding labels in  $\{0, \dots, C - 1\}$ , and let  $P(Y = k|X_i = x_i, \theta)$  the probability for class  $k$  given input sample  $x_i$ , then the cross-entropy loss for one sample  $(x_i, y_i)$  is defined as:

$$\sum_{k=0}^{C-1} -\tilde{y}_i[k] \ln P(Y = k|X_i = x_i, \theta)$$

- This is again the neg-log probability for the class of the ground-truth label
  - if the probability for the ground-truth class is  $= 1$ , we have  $-\ln(1) = 0$  for the loss

## Takeaway for today:

- logistic regression: model which computes inner product  $+b$  with one output, then combines it with the logistic sigmoid
- a possible loss: binary cross entropy - the neg-log of the output probability for the ground truth class  $y$  for a given sample  $(x, y)$
- neg-log (close to 0) = close to  $\infty$ ,  $-\ln(1.0) = 0$
- binary cross entropy - derived from maximum likelihood principle to observe the ground truth labels
- logistic regression: 1-layer NN with sigmoid activation function

The key you need to understand sigmoid activation and cross entropy loss, is how the exponential and the (neg) logarithm look like.